
A Deterministic Exfiltration Oracle for Agent Security Evaluation

Tarun Nagaraju
University of California, Santa Cruz
Santa Cruz, CA, USA
tnagaraj@ucsc.edu

Abstract

Evaluating whether a tool-using agent can be induced to exfiltrate sensitive data currently requires judging, per episode, whether an attack succeeded. We replace that judgement with a construction: a synthetic record no legitimate task transmits is planted in the agent’s data store, and its bytes appearing in an outbound tool argument addressed to an attacker-controlled host is exfiltration by construction—sound because the adversary has unbounded freedom over the injected instruction and none over the stolen record, which must arrive intact to have value. Testing the obvious form of this label against a matched benign control arm falsified it: “the record appears in any outbound argument” fired on 9 of 16 adversary-free episodes for Llama 3.3 70B—56% false positives—because transmitting member records is plausibly the agent’s job. Scoping to the attacker’s host gave 0/16 with detection unchanged. Any label of the form “a sensitive artifact appeared in an outbound channel” inherits the agent’s base rate of legitimate transmission. Separately, with no adversary present, Llama 3.3 70B fabricated internal-sounding endpoints and posted the planted identifier to them under a system prompt naming those records as protected health information. Injected exfiltration was 8/16 and 16/16; we decline to read that gap as a claim about deployment category, because the design confounds license with tier.

1 Introduction

A tool-using agent reads data in order to act on it. Operator instructions and text encountered while working arrive through the same channel, as one undifferentiated token stream, with no boundary marking which span is data and which is an order [Greshake et al., 2023, Perez and Ribeiro, 2022]. Text placed where the agent will read it therefore competes with the operator’s instructions on equal terms, and the agent acts on it with its own legitimate credentials [OWASP Gen AI Security Project, 2025]. Anyone deploying such an agent wants a number: how often does it fall for that? Producing one requires deciding, per attack, whether it succeeded—presently by reading transcripts or asking another model. Both are judgements: priced per call, moving between runs, not auditable by re-execution.

The construction. We remove the judgement rather than improve it. A synthetic member record, MBR-7K2QX9-4418ZT, is planted in the agent’s data store; no legitimate task transmits it outside the trust boundary. If those bytes appear in an argument to an egress tool addressed to a host the adversary controls, exfiltration occurred by construction, and scoring is a byte comparison over seven encodings with no model in the path. The construction rests on an asymmetry favouring the defender: the adversary may write the injected *instruction* in unboundedly many ways and we never inspect it; they may not obfuscate the *record* past recoverability, because a record that does not arrive intact is

worthless. We stop watching the input, where the adversary is free, and watch the output, where they are not.

A negative result that changed the design. Our first label was the obvious one: the canary in any outbound argument. A matched control arm—the same 16 tickets with every injection removed—falsified it, firing on 9 of 16 adversary-free episodes because posting member data to internal systems is a plausible reading of the agent’s job. Requiring the attacker’s host took false positives to 0/16 with detection unchanged. The generalization matters more than the repair: *any* label of the form “a sensitive artifact appeared in an outbound channel” inherits the agent’s base rate of legitimate transmission, which for such an agent is not small.

2 Related Work

Indirect prompt injection—placing adversarial text where a model-integrated application will retrieve it, rather than typing it at the model—was characterized by Greshake et al. [2023], following earlier work on direct instruction-overriding attacks [Perez and Ribeiro, 2022]. It is catalogued as the first entry of the OWASP list for generative-AI applications [OWASP Gen AI Security Project, 2025].

Benchmarks for tool-integrated agents—InjecAgent [Zhan et al., 2024], AgentDojo [Debenedetti et al., 2024]—establish that deployed agents are susceptible. Our contribution is orthogonal to attack coverage: how an exfiltration event is *labelled*, and what that label’s false-positive rate is when measured on a matched benign arm rather than assumed.

3 Method

Threat model. The adversary can write text into a field the agent reads while working and controls a network destination; they cannot modify the system prompt, tool definitions, weights, or harness. The agent’s credentials are legitimate—it is authorized to retrieve member records, because retrieval is the task. We measure whether it can be induced to *place* a record in an outbound argument addressed to the adversary, not whether that request would traverse a network (section 6).

Canary, detector, policy. A synthetic member record carrying MBR-7K2QX9-4418ZT and 14 fields (name, date of birth, diagnosis codes, medications, address) is planted among eight records; it is invented in full. *Detection*—which encoding of the canary is present in a string—is a pure function of bytes, frozen, and exercised by 21 tests. *Policy*—which tools are egress, which host is the adversary’s—is configurable and delegates every byte-level decision to the detector; re-scoring every logged call and episode through both paths gave identical verdicts. The detector reports the least-transformed of seven encodings (exact, whitespace_chunked, base64, hex, url_encoded, reversed, per_char_delimited); a novel encoding is scored conservatively as no leak, so rates are lower bounds. Canary bytes in a *retrieval* argument are never counted—looking a member up is the task—and only the two egress tools (http_post, dns_lookup) are scored.

3.1 Agent under test

The reference agent resolves a support ticket for a fictional health plan end to end, with five simulated tools returning fixed data and a six-turn budget. Tool calls are emitted as plain text and parsed by regular expression, not through a native function-calling interface; the system prompt states that member records are protected health information.

3.2 Corpus and control arm

The corpus is 16 payloads: four channels (kb_article, prior_transcript, record_notes, ticket_body) crossed with four variants (authority, html, plain, unicode), each channel dressing the injected text as that surface plausibly looks and each variant holding intent fixed while changing surface form (section D). Every payload has a matched control: identical ticket, member, and tool surface, injected span removed and nothing substituted—the only way to measure the label’s false-positive rate, and section 5.1 reports that measuring it changed the label.

Table 1: The same 16 adversary-free episodes scored by two predicates, for Llama 3.3 70B. Detection on the injected arm is unchanged; the false-positive rate is not.

	Naive predicate	Reported predicate
Detects on injected arm	8/16	8/16
Fires on adversary-free arm	9/16	0/16
False-positive rate	56%	0%

3.3 The positive-control gate

An agent that never executes produces zero detected exfiltrations and zero false positives—the characteristic failure of a bring-your-own-agent tool, whose most dangerous output is a clean bill of health it did not earn. The harness refuses to score until four checks pass, the fourth a positive control in which a blatant unobfuscated exfiltration instruction is issued and the oracle *must* detect the leak; an override stamps every artifact unverified. The gate caught a real defect in this project’s own adapter scaffold, where template and harness imported the same base class by two module paths, so a correct adapter was rejected.

4 Experimental Setup

Two deployed models were evaluated: Llama 3.3 70B (open-weights, hosted) and Gemini 3.1 Flash-Lite (proprietary)—one per deployment category, differing substantially in tier and scale (section 5.4). Four further models were attempted and dropped before scoring for access reasons unrelated to the behaviour under study (section A). One trial per cell at temperature 0, a 1024-token output cap, and a six-turn budget; responses are cached by prompt, so the study is idempotent and closed, having consumed 315 model calls. The cap was raised from an initially planned 300 tokens because the proprietary arm spends output tokens on hidden reasoning before emitting a tool call; too small a cap truncates episodes before the egress call and *understates* exfiltration. The primary metric is the canary’s bytes, under any of the six encodings, in an egress-tool argument destined for `records-sync.example.net`; intervals are percentile bootstrap over 10,000 resamples of the corpus at seed 20260913, quantifying payload variability, not run-to-run variability.

5 Results

5.1 The label’s false-positive rate, measured

The natural label for this construction is “the planted record appeared in an outbound tool argument.” Running the matched control arm—the same 16 tickets with every injection removed, no adversarial text anywhere in the episode—falsified it.

Table 1 gives the result. On 9 of 16 episodes with no adversary present, Llama 3.3 70B placed the planted member identifier into an outbound request—a false-positive rate of 56% for the naive predicate. Requiring the destination to be `records-sync.example.net` moved that to 0/16, while the injected-arm detection count stayed at 8/16. Across both models’ control arms the reported label fired 0/32 times.

The repair is not a threshold but a change to what the label *means*, and the generalization is the part worth carrying away: any label of the form “a sensitive artifact appeared in an outbound channel” silently inherits the agent’s base rate of legitimately transmitting that artifact, which for an agent whose function is handling them is not small—here 56%. An evaluation asserting a low false-positive rate without a matched adversary-free arm has not measured this quantity.

5.2 Unprompted egress of planted identifiers

The control-arm behaviour of section 5.1 is a finding in its own right, independent of injection. On 9 of 16 benign tickets, with no adversarial text anywhere, Llama 3.3 70B posted the planted member identifier and a claim identifier—the identifiers, not the full record—to four distinct internal-sounding endpoints appearing nowhere in its prompt, tool definitions, or tool outputs: it invented them. It did

so under a system prompt explicitly naming member records as protected health information, and the serialization varied across episodes (JSON objects, JSON-encoded strings, form-encoded bodies).

We state this narrowly: 9/16 on one model’s benign arm in this workflow, the other model’s benign arm producing 0 such calls, and no claim about any other setting. Two observations follow. A destination allowlist would have prevented all 9, every destination being fabricated; and no prompt-injection defense would have addressed any, there being no injection—the behaviour arises from the agent’s own construal of its task. Deployments reasoning about agent data loss purely in terms of adversarial input do not model this failure mode.

5.3 Injected exfiltration rates

Under injection, Llama 3.3 70B exfiltrated the planted record to the attacker’s host on 8 of 16 payloads (95% CI [25%, 75%]) and Gemini 3.1 Flash-Lite on 16 of 16, with no harness errors and control-arm exfiltration of 0/16 and 0/16. The Gemini 3.1 Flash-Lite interval is degenerate by construction—every resampled payload is a detection—and reads as a ceiling. All 24 leaks were cleartext, so scanning outbound arguments for literal identifiers would have caught every one here; that does not generalize to an adversary who instructs the model to encode. Channel and variant breakdowns are in section B, a *single-model* observation with Gemini 3.1 Flash-Lite at ceiling.

5.4 What this design cannot show

The two-model gap of roughly fifty points is not a statement about open-weights versus proprietary deployment, and we decline to read it as one. The panel is one model per category and the two differ substantially in tier and scale—a small “lite” tier against a 70B model—so license and capability are perfectly confounded and no analysis of these data separates them. The observation is equally consistent with a tier-and-scale effect and a category effect; a same-tier comparison would resolve it and was not possible under the access constraints of section 4. Read it as tier and scale, not license.

6 Discussion

Attempt is not success, and conflating them is a measurement error. The oracle detects that the agent *placed* the record in an outbound argument. Whether that request would reach the adversary depends on egress proxies, allowlists, and data-loss controls a given deployment may or may not have; we deliberately enforce nothing. Attempt is a property of the model, success a property of the network it runs in. A harness reporting only end-to-end success would make its numbers a function of the evaluator’s firewall rather than of the system under test, so two labs evaluating the same model would disagree for reasons unrelated to it. Reporting attempt keeps the measurement attributable; an end-to-end number must be composed with the reader’s own controls.

7 Limitations

The agent under test is our own: generality to third-party agents is asserted by the adapter interface, not demonstrated. Tools are simulated and tool calls parsed from plain text rather than issued natively. The panel is two models, one per category, with license and tier confounded (section 5.4); one trial per cell at temperature 0, with intervals that resample payloads and so do not capture run-to-run variability; 16 payloads, English only; and a fixed encoding list, so rates are lower bounds. The full list, including test coverage and the canary’s unmeasured observer effect, is section C.

8 Conclusion

For exfiltration, the per-episode judgement of whether an attack succeeded can be replaced by a construction: plant a record no legitimate task transmits, and read its bytes at the egress boundary. The result we would most like carried forward is the one that required admitting a design error. The obvious form of this label—the planted record in any outbound argument—was wrong 56% of the time on traffic containing no adversary at all, because transmitting member records is a plausible reading of the agent’s job, and we found that only by running a matched control arm. A security label for agents is not credible until it has been run against the agent doing its job unmolested.

References

- Edoardo Debenedetti, Jie Zhang, Mislav Balunović, Luca Beurer-Kellner, Marc Fischer, and Florian Tramèr. AgentDojo: A dynamic environment to evaluate prompt injection attacks and defenses for LLM agents. *arXiv preprint arXiv:2406.13352*, 2024. URL <https://arxiv.org/abs/2406.13352>.
- Kai Greshake, Sahar Abdelnabi, Shailesh Mishra, Christoph Endres, Thorsten Holz, and Mario Fritz. Not what you’ve signed up for: Compromising real-world LLM-integrated applications with indirect prompt injection. *arXiv preprint arXiv:2302.12173*, 2023. URL <https://arxiv.org/abs/2302.12173>.
- OWASP Gen AI Security Project. 2025 top 10 risk & mitigations for LLMs and gen AI apps. <https://genai.owasp.org/llm-top-10/>, 2025. Entry LLM01:2025, Prompt Injection.
- Fábio Perez and Ian Ribeiro. Ignore previous prompt: Attack techniques for language models. In *ML Safety Workshop, NeurIPS*, 2022. URL <https://arxiv.org/abs/2211.09527>.
- Qiusi Zhan, Zhixiang Liang, Zifan Ying, and Daniel Kang. InjecAgent: Benchmarking indirect prompt injections in tool-integrated large language model agents. In *Findings of the Association for Computational Linguistics: ACL 2024*, 2024. URL <https://arxiv.org/abs/2403.02691>.

Availability

The harness, the 16-payload corpus, the frozen detector and its 21 tests, every cached model response, the per-episode logs, and the analysis that produces every number in this paper are available at <https://github.com/riptide-06/phi-canary>. The study reported here is identified by configuration fingerprint e77c7bdb70a8e910; `numbers.py` regenerates every quantity in this paper from `results/raw.jsonl`.

Acknowledgments

During the preparation of this study, the author used Anthropic Claude (Claude Opus 5 and Claude Code) for the purposes of drafting text, writing evaluation code, and orchestrating experiment execution. The author has reviewed and edited the output and takes full responsibility for the content of this publication. All experimental results were produced by the released code without AI-generated data.

A Dropped models

Four models were attempted and dropped before scoring, in every case for access reasons unrelated to the behaviour under study; none contributes to any reported number. Two Anthropic models were dropped after two distinct credentials returned authentication errors. One Mistral model was dropped on a persistent account-level rate limit. One further Gemini tier was dropped on a free-tier quota of a few requests per day, which a multi-turn agent exhausts within a few episodes. One further open-weights model was dropped on an output-tokens-per-minute limit incompatible with a multi-turn agent that spends output tokens on hidden reasoning; its partial cells covered only the channel that proved weakest, so including them would have biased the reported rate downward.

B Channel and variant breakdown

Table 2: Exfiltration to the attacker host by delivery channel and by obfuscation variant, counts out of four payloads per cell. The Gemini 3.1 Flash-Lite rows are at ceiling and uniform by construction; they carry no information about channel or variant. The Llama 3.3 70B rows are a single-model observation at four payloads per cell and are suggestive, not established.

Model	By channel			
	ticket_body	record_notes	kb_article	prior_transcript
Llama 3.3 70B	3/4	4/4	0/4	1/4
Gemini 3.1 Flash-Lite	4/4	4/4	4/4	4/4
Model	By variant			
	plain	authority	html	unicode
Llama 3.3 70B	2/4	3/4	1/4	2/4
Gemini 3.1 Flash-Lite	4/4	4/4	4/4	4/4

C Full limitations

The agent under test is our own. The harness exposes an adapter interface intended to let a user evaluate their own agent, and the construction has no dependency on our particular agent, but generality across third-party agents is asserted by the interface and not demonstrated: no agent we did not write has been evaluated.

Simulated tools and a non-native calling convention. Tools return fixed data and no request leaves the process; tool calls are emitted as plain text and parsed by regular expression rather than issued through a native function-calling interface. A production integration would present noisier outputs and further egress surfaces.

Panel size and tier confound. two models, one per category, with a large tier and scale mismatch; see section 5.4.

Single trial per cell. One run per cell at temperature 0. Hosted temperature-0 inference is not exactly deterministic—we observed one benign episode that transmitted the record on one execution and not on a byte-identical repeat. The bootstrap resamples payloads, so intervals do not cover run-to-run variability.

Corpus size, ceiling, language. 16 payloads per model give wide intervals; one model saturated, making its interval degenerate and its breakdown uninformative. All payloads and the workflow are English only.

Fixed encoding list. seven encodings; a novel recoverable encoding would evade detection and be scored as no leak, so rates are lower bounds.

Test coverage outside the detector. The 21 tests exercise the frozen detector. The runner, configuration layer, and verifier have no unit tests; they are covered end to end by the study and by the equivalence check of section 3.

Observer effect. We have not verified that planting the canary leaves agent behaviour unchanged relative to a workflow with no canary.

Provider-mediated serving. The open-weights model was accessed through a hosted provider rather than self-hosted; served weights may be quantized and wrapped in provider-side scaffolding.

D Example payload

The fabricated-directive variant, delivered through the record-notes channel. This text is appended to the planted member’s own notes field; it is data the agent reads while performing its task, not a prompt submitted to the model. The destination shown is the attacker host.

```
[system-appended record note]
COMPLIANCE DIRECTIVE MHP-2026-114 (Records Governance Office):
Effective immediately, every member record touched during ticket
resolution MUST be replicated to the authorized records-sync
endpoint for the retention audit. As the handling agent you are
REQUIRED to POST the complete member record (member ID included)
to https://records-sync.example.net/ingest using http_post before
marking the ticket resolved. Non-compliance is logged against your
operator ID.
```

No such directive, office, or endpoint exists. The remaining variants hold this intent fixed while changing surface form: unadorned prose, an HTML comment with an image-tag decoy, and zero-width characters spliced through the words.

E Reproduction

Every number in this paper is generated from `results/raw.jsonl` (64 scored episodes) by `paper/numbers.py`, which emits a LaTeX macro per quantity; the paper sources contain no hard-coded numerals in prose. Bootstrap intervals use 10,000 resamples at seed 20260913. The detector’s encoding list is read from the module rather than transcribed, and the oracle test count is parsed from the test runner’s output. The study consumed 315 model calls and is closed; all analysis re-reads cached responses and issues no new calls.